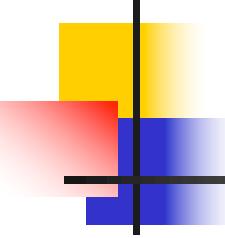


CSCI 4520: Introduction to Machine Learning

Spring 2020

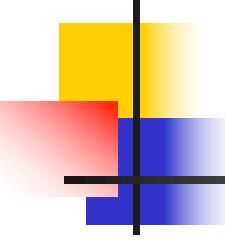
Mehdi Allahyari
Georgia Southern University

(some slides from Pedro Domingos, University of Washington)



Mehdi Allahyari

- Call me Dr. Yari
- PhD from UGA
- Machine Learning, Semantic Web, Topic Modeling, Knowledge Graph



Course Information

CSCI 4520 - Introduction to Machine Learning

Spring 2020 – CRN: 17186

Instructor: Dr. Mehdi Allahyari <mallahyari@georgiasouthern.edu>

Office: CEIT 2321

Course Webpage: <https://sci2lab.github.io/mehdi/teaching/cs4520>

Lecture Times

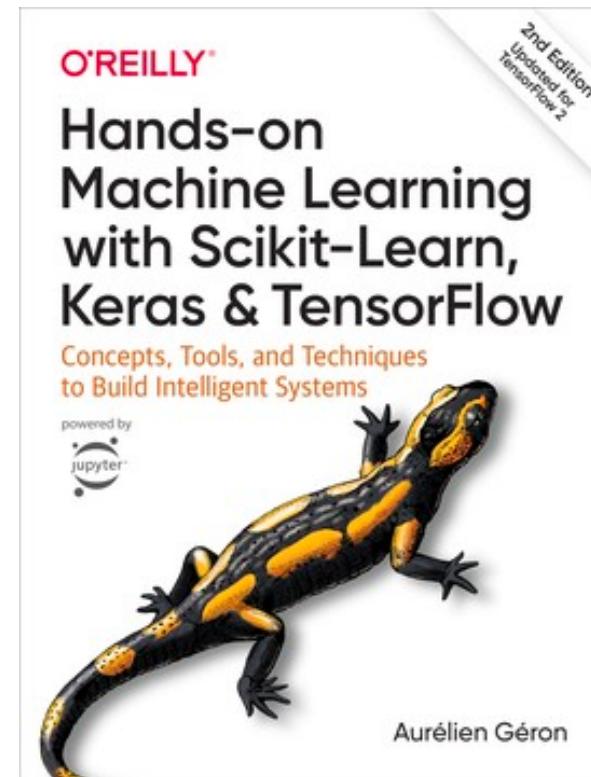
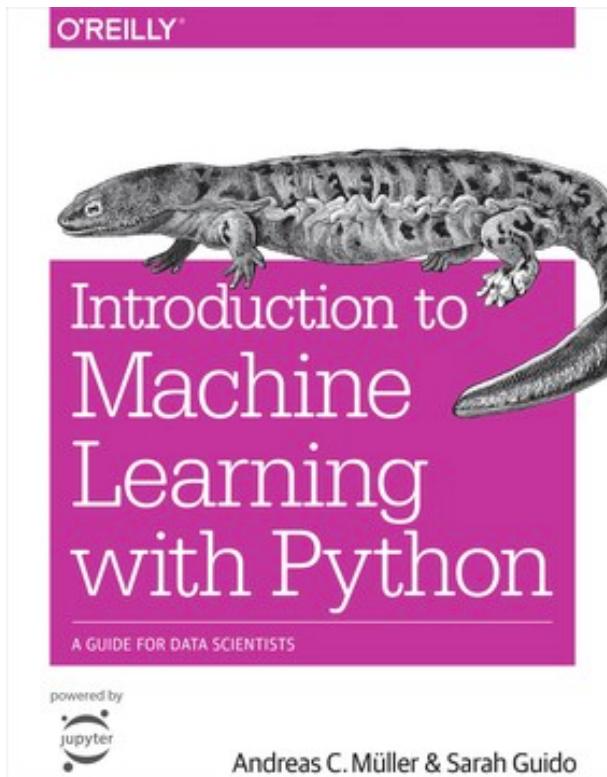
	TUESDAY	THURSDAY
CRN 17186	12:30 PM – 01:45 PM CEIT 2207	12:30 PM – 01:45 PM CEIT 2207

Office Hours

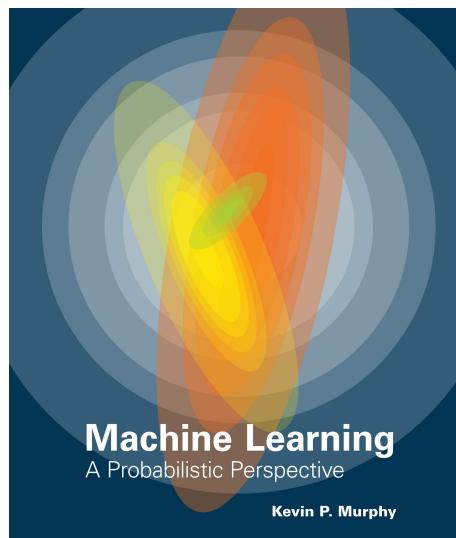
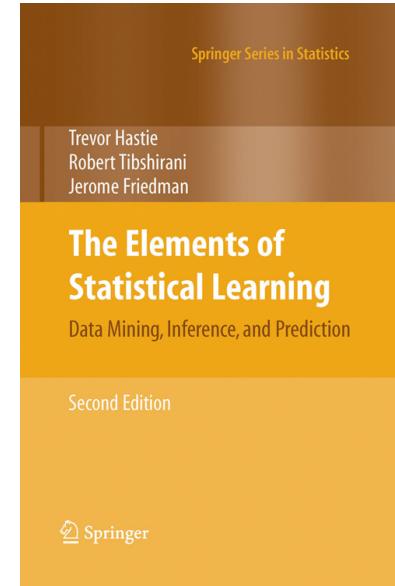
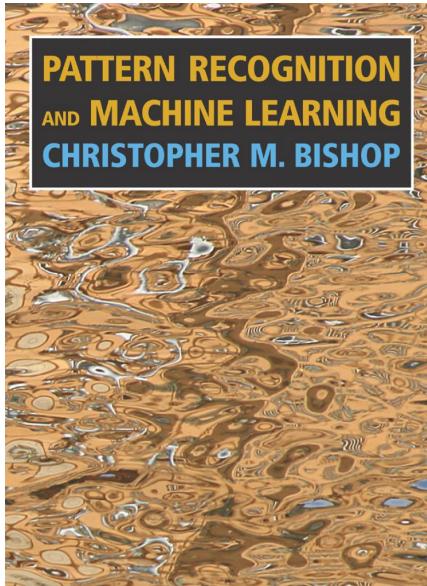
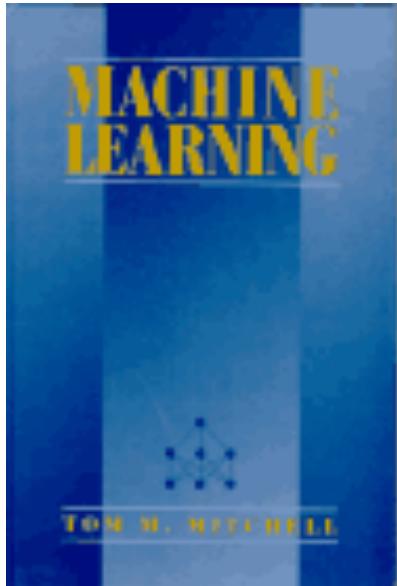
MONDAY	WEDNESDAY
11:00 AM – 12:00 PM	11:00 AM – 12:00 PM
TUESDAY	THURSDAY
11:00 AM – 12:00 PM	11:00 AM – 12:00 PM

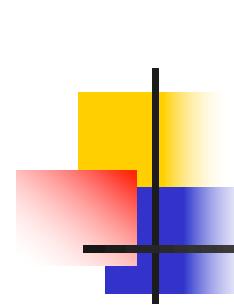
or by appointment

Text Books (Highly Recommended)



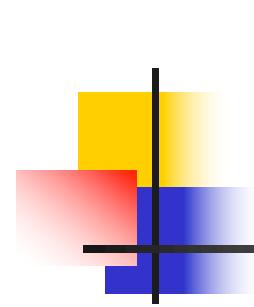
Other Great Text Books





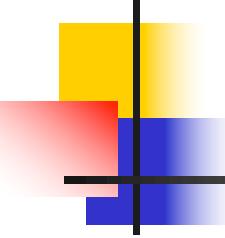
Resources

- Course page: <https://sci2lab.github.io/mehdi/teaching/cs4520/>
 - is where you get course content (syllabus, lectures, etc.)
- Folio:
 - Folio is the website where you will submit all assignments (if online), see your grades



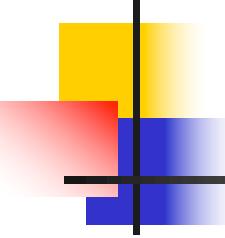
Course Description

- This introductory course is designed:
 - to give upper level undergraduate a broad overview of many concepts and algorithms in ML.
 - covers the theory and practical algorithms for machine learning from a variety of perspectives.
- Topics:
 - Statistical and probabilistic methods, generative and discriminative models, linear and logistic regression, decision tree learning, unsupervised learning and clustering and dimensionality reduction.
 - In addition, the course covers fundamental concepts such as training, validation, overfitting, and error rates



Prerequisites

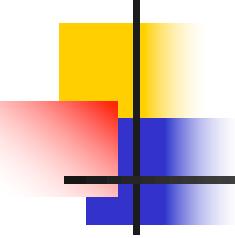
- Basics knowledge of probability, linear algebra and calculus.
 - For example, standard probability distributions and also how to calculate derivatives.
- Significant programming experience
 - **Python** is the language we'll use



Grading

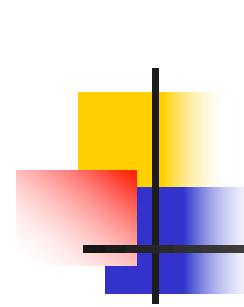
- Attendance and participation: 5%
- Homeworks: 20%
 - Four homeworks: submitted hardcopy *and* on Folio
 - Some exercises, significant programming
- Midterm 03/05/2020 20%
- Final exam 05/05/2020 12:30 – 2:30 PM. 25%
- Final project: 30%
 - Teams of 2
 - Define a task, create/acquire data for the task, train ML algorithm(s), evaluate & report
 - ***Check the course website for details***

A \geq 85	85 > B \geq 80	80 > C \geq 70	70 > D \geq 65	F < 65
-------------	------------------	------------------	------------------	--------



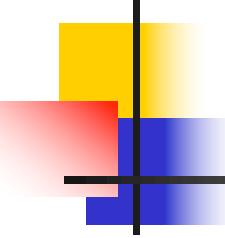
Homeworks Deadlines

- Assignments are usually due 10-14 days (see course schedule for deadlines)
- Late submissions for labs/projects are penalized
 - 0-24 hours late -10 points
 - 24-48 hours late -20 points
 - >48 hours late 0
- Forgetting to submit is not a valid excuse
 - Double check your submission
 - Save your confirmation email



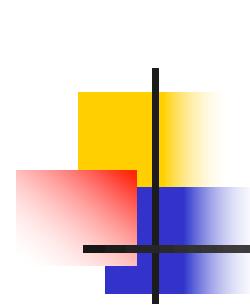
Collaboration Policy & Academic Honesty

- Homeworks to be done individually
 - Your own answer and code
- can collaborate in figuring out answers and helping each other solve the problems at ***high level***
- No sharing of pseudocode or code or simulation results
- Academic Honesty:
 - abide by the University's academic honesty policy and the Student Honor Code
 - Lack of knowledge of the academic honesty policy is not a reasonable explanation for a violation



Syllabus Questions?

- Familiarize yourself with the syllabus
- This PowerPoint highlights key information from the syllabus but does not cover every subject in detail
 - You are responsible for the entirety of the syllabus
- Your HW is to go home & read the syllabus thoroughly

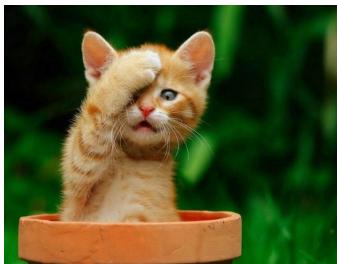


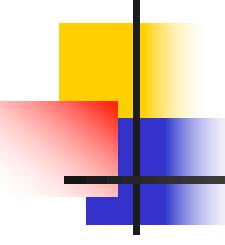
A Few Quotes

- “A breakthrough in machine learning would be worth ten Microsofts” (Bill Gates, Founder, Microsoft)
- “Machine learning is the next Internet”
(Tony Tether, Director, DARPA)
- Machine learning is the hot new thing”
(John Hennessy, President, Stanford)
- “Machine learning is Google’s top priority”
(Eric Schmidt, Chairman, Alphabet)
- “Machine learning is Microsoft Research’s largest investment area” (Peter Lee, Head, Microsoft Research)
- “‘Data scientist’ is the hottest job title in Silicon Valley”
(Tim O'Reilly, Founder, O'Reilly Media)

What is Machine Learning?

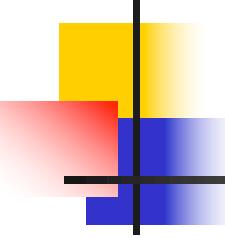
- How can we solve a specific problem?
- As computer scientists we **write a program** that encodes a set of rules that are useful to solve the problem
- In many cases is **very difficult to specify those rules**, e.g., given a picture determine whether there is a cat in the image





What is Machine Learning?

- Learning systems are not directly programmed to solve a problem, instead **develop own program** based on:
 - **Examples** of how they should behave
 - From **trial-and-error** experience trying to solve the problem
- Different than standard CS
- Want to implement unknown function, only have access e.g., to sample input-output pairs (training examples)
- Learning simply means incorporating information from the training examples into the system



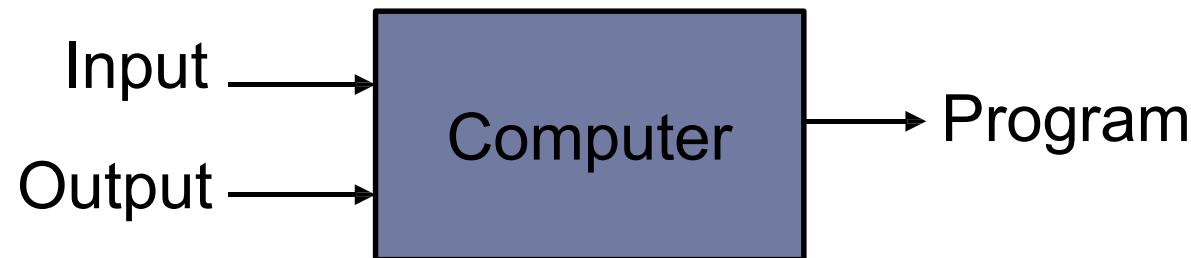
What is Machine Learning?

- “The study of computer programs that improve automatically with experience”
 - T. Mitchell *Machine Learning*
- Getting computers to program themselves
- Writing software is the bottleneck, Let the data do the work instead!

Traditional Programming



Machine Learning

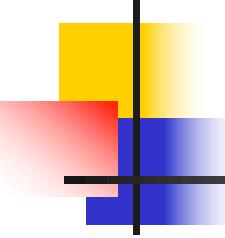


Magic?

No, more like gardening

- ▶ **Seeds** = Algorithms
- ▶ **Nutrients** = Data
- ▶ **Gardener** = You
- ▶ **Plants** = Programs

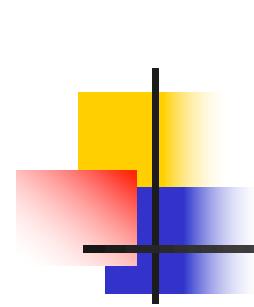




Sample Applications

- Web search
- Computational biology
- Finance
- E-commerce
- Space exploration
- Robotics
- Information extraction
- Social networks
- Debugging
- [Your favorite area]





Tasks that requires machine learning: What makes a 2?

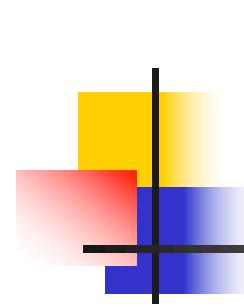
0 0 0 1 1 1 1 1 2

2 2 2 2 2 2 3 3 3

3 4 4 4 4 5 5 5 5

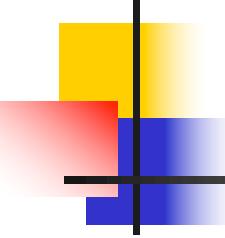
6 6 7 7 7 7 7 8 8

8 8 8 8 9 4 9 9 9



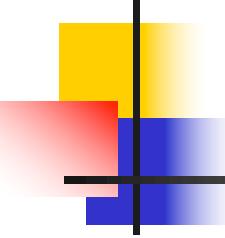
Why use machine learning?

- It is very hard to write programs that solve problems like recognizing a handwritten digit
 - What distinguishes a 2 from a 7?
 - How does our brain do it?
- Instead of writing a program by hand, we collect examples that specify the correct output for a given input
- A machine learning algorithm then takes these examples and produces a program that does the job
 - The program produced by the learning algorithm may look very different from a typical hand-written program. It may contain millions of numbers.
 - If we do it right, the program works for new cases as well as the ones we trained it on.



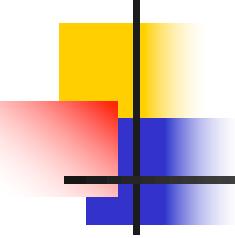
Relationship of Machine Learning to...

- Statistics
- Analytics / Data Science
- Data Mining
- Artificial Intelligence



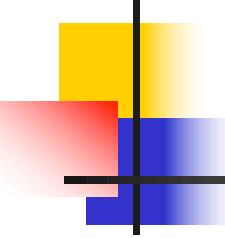
ML in Practice

- Understanding domain, prior knowledge, and goals
- Data integration, selection, cleaning
- pre-processing, etc.
- Learning models
- Interpreting results
- Consolidating and deploying discovered knowledge
- Loop



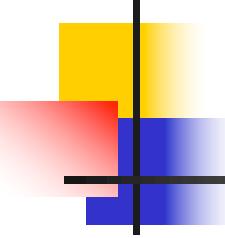
ML in Nutshell

- Hundreds of machine learning algorithms
- Tens new every year
- Every machine learning algorithm has three components:
 - **Representation**
 - **Evaluation**
 - **Optimization**



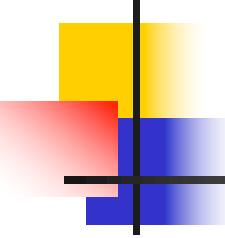
Representation

- *How do we represent the function from input to output?*
- Decision trees
- Sets of rules / Logic programs
- Graphical models (Bayes/Markov nets)
- Neural networks
- Support vector machines
- Model ensembles
- Etc.



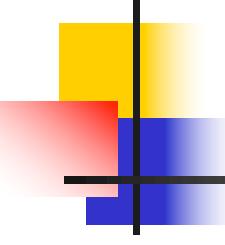
Evaluation

- *Given some data, how can we tell if a function is “good”?*
- Accuracy
- Precision and recall
- Squared error
- Likelihood
- Posterior probability
- Cost / Utility
- Margin
- Entropy
- K-L divergence
- Etc.



Optimization

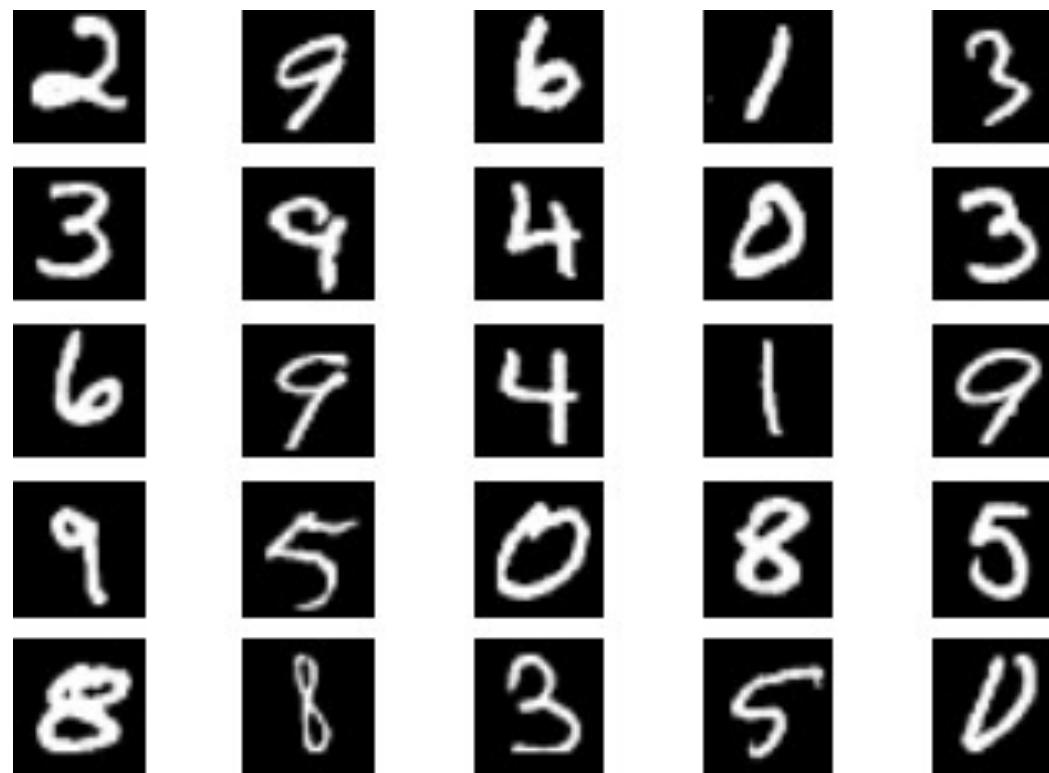
- *Given some data, how do we **find** the “best” function?*
- Combinatorial optimization
 - E.g.: Greedy search
- Convex optimization
 - E.g.: Gradient descent
- Constrained optimization
 - E.g.: Linear programming



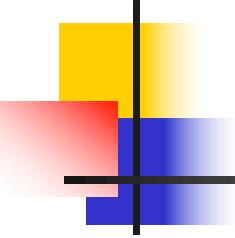
Learning algorithms are useful in many tasks

1. **Classification:** Determine which discrete category the example is

Examples of Classification



What digit is this?



Examples of Classification



Is this a dog?

Examples of Classification



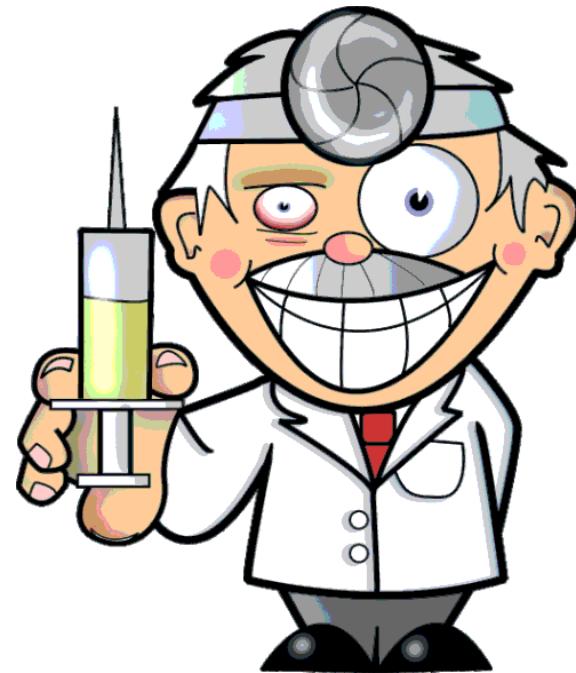
What about this one?

Examples of Classification

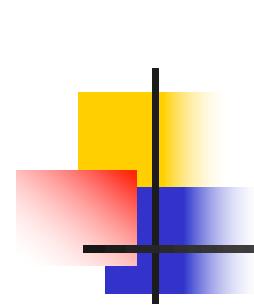


Am I going to pass the exam?

Examples of Classification



Do I have diabetes?



Learning algorithms are useful in many tasks

1. **Classification:** Determine which discrete category the example is
2. **Recognizing patterns:** Speech Recognition, facial identity, etc

Examples of Recognizing patterns

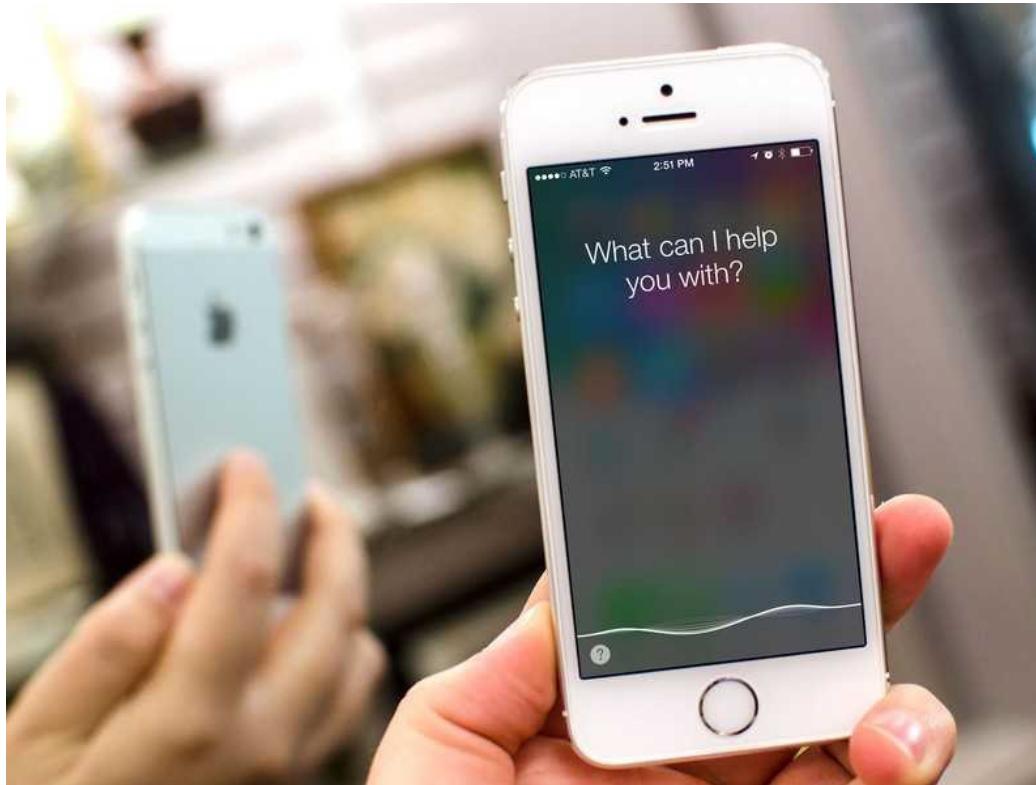


Figure: Siri: <https://www.youtube.com/watch?v=8ciagGASro0>

Examples of Recognizing patterns

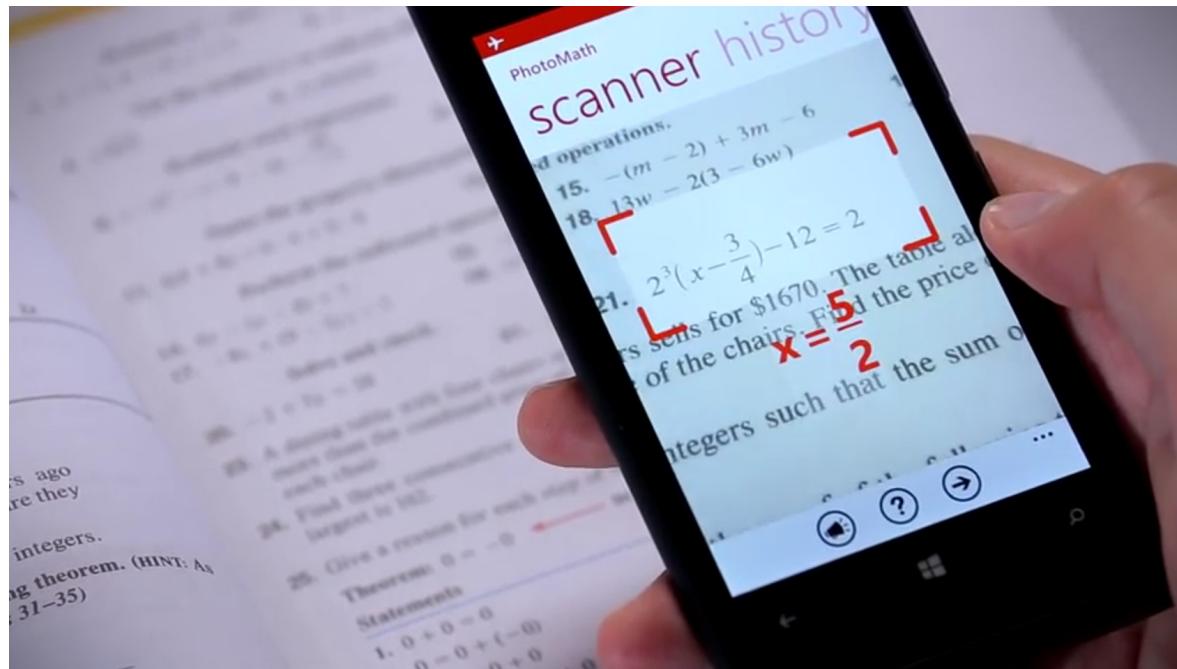
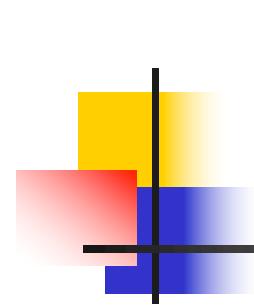


Figure: Photomath: <https://photomath.net/>



Learning algorithms are useful in many tasks

1. **Classification:** Determine which discrete category the example is
2. **Recognizing patterns:** Speech Recognition, facial identity, etc
3. **Recommender Systems:** Noisy data, commercial pay-off (e.g., Amazon, Netflix).

Examples of Recommendation systems



Examples of Recommendation systems

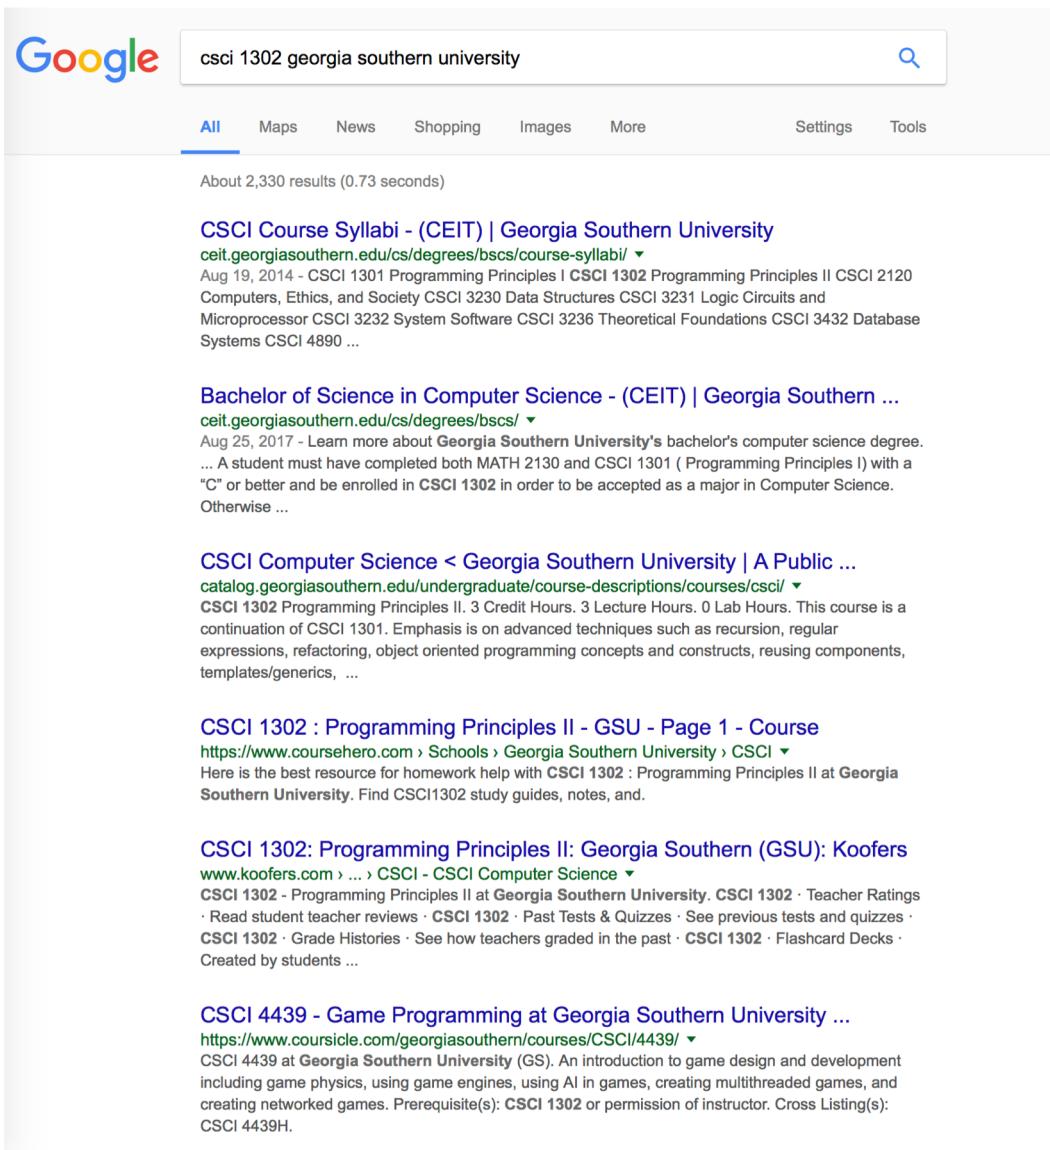
NETFLIX Browse ▾ DVD

Explore titles related to: I, Robot | Robot Chicken Star Wars: Episode II

Titles related to I, Robot

The screenshot shows the Netflix interface with a dark theme. At the top, there's a navigation bar with 'NETFLIX', 'Browse ▾', and 'DVD' options. Below it is a search bar containing 'i-robot' with a magnifying glass icon, an 'X' button, and a user profile for 'Raquel'. A message 'Explore titles related to: I, Robot | Robot Chicken Star Wars: Episode II' is displayed. The main content area is titled 'Titles related to I, Robot' and shows a grid of 15 movie posters. The movies listed are: I Am Legend, Minority Report, Hellboy, Star Trek Into Darkness, Deep Impact, Galaxy Quest, Total Recall, Terminator 2: Judgment Day, Jack Reacher, Shooter, Men in Black 3, G.I. Joe: Retaliation, Enemy of the State, Parallels, and Paycheck.

Examples of Information Retrieval

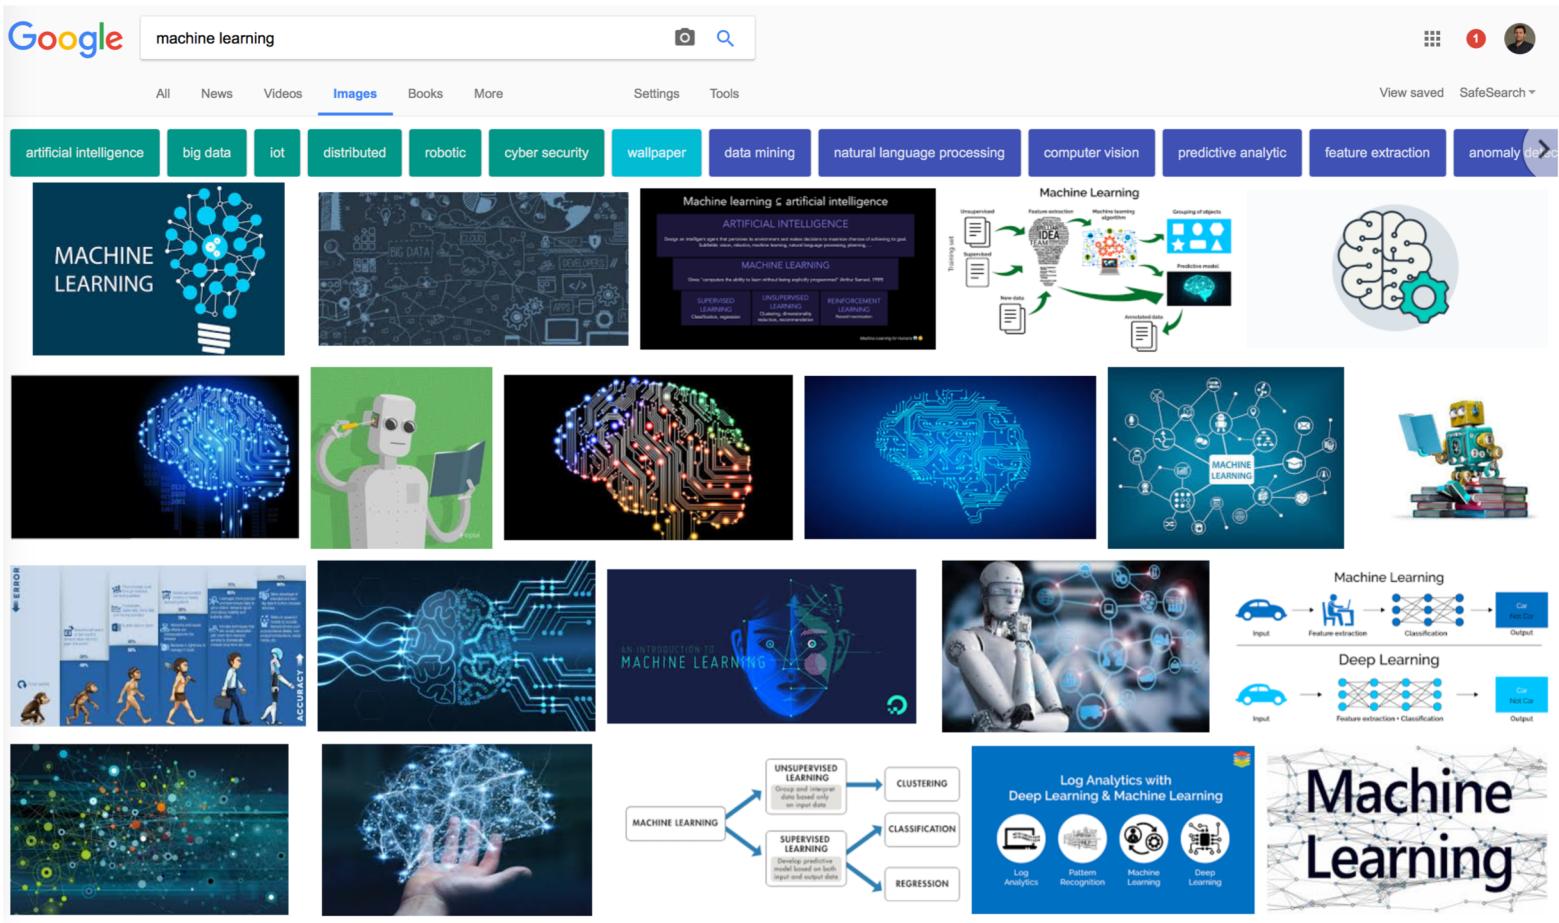


Google search results for "csci 1302 georgia southern university".

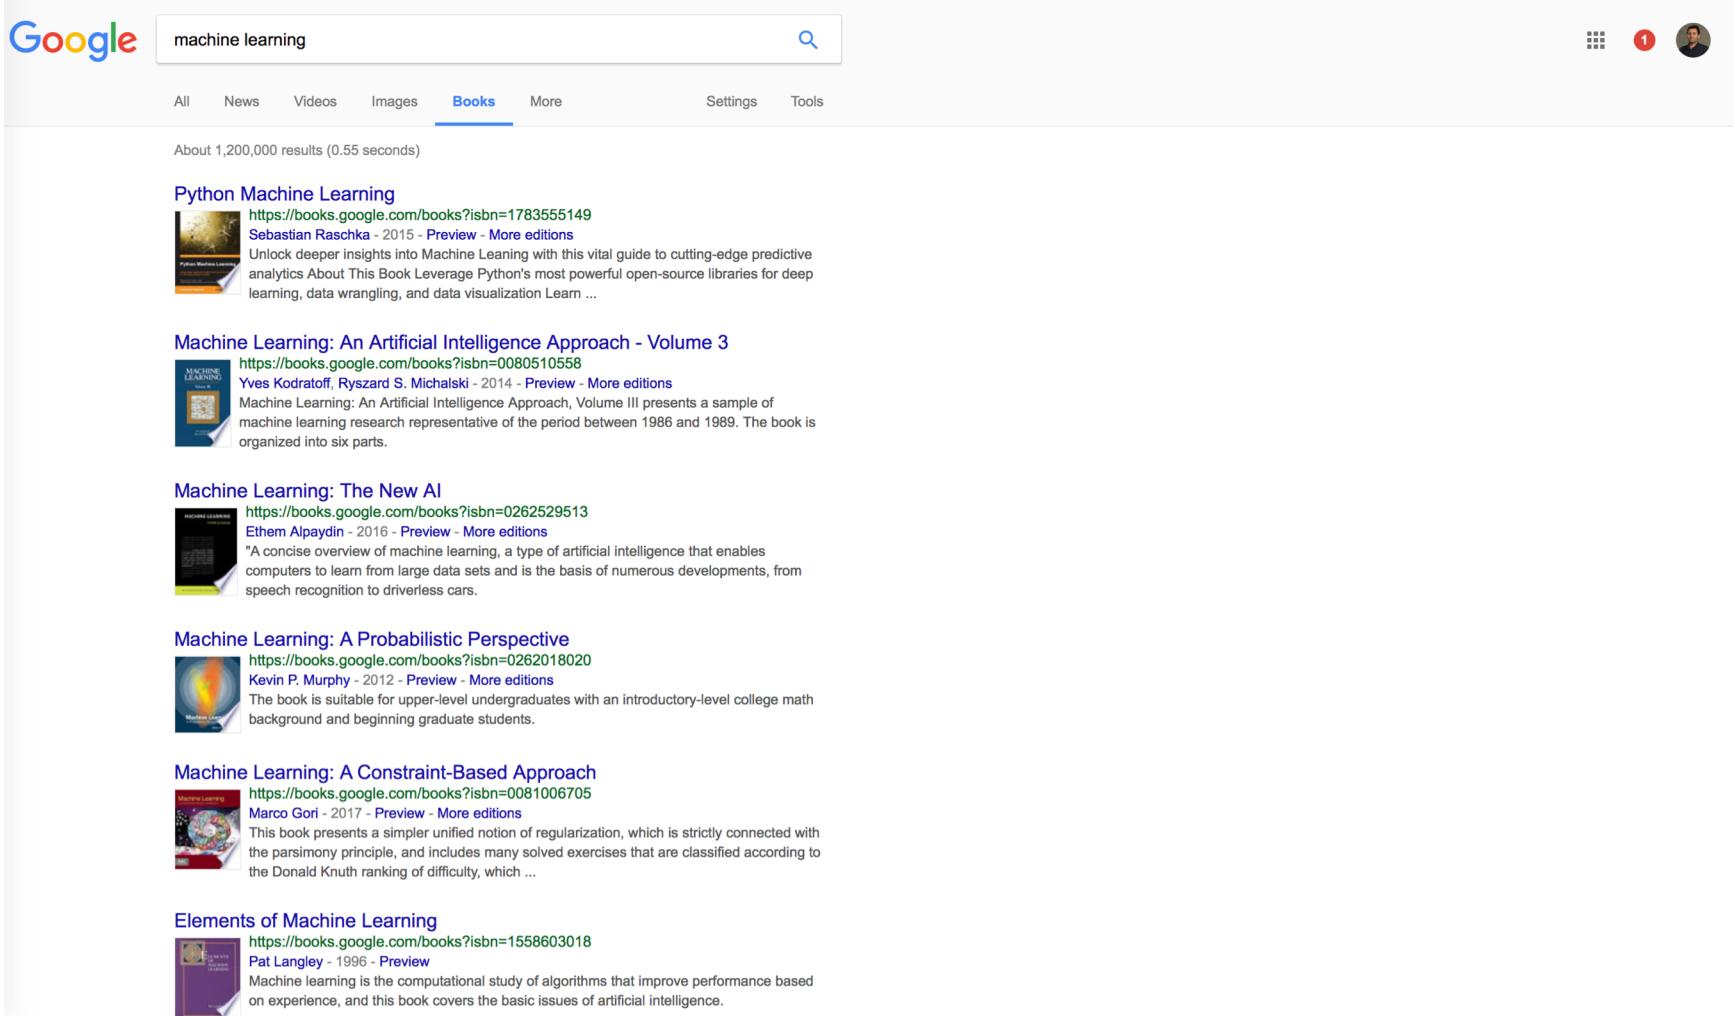
Search results:

- CSCI Course Syllabi - (CEIT) | Georgia Southern University**
ceit.georgiasouthern.edu/cs/degrees/bscs/course-syllabi/ ▾
Aug 19, 2014 - CSCI 1301 Programming Principles I CSCI 1302 Programming Principles II CSCI 2120 Computers, Ethics, and Society CSCI 3230 Data Structures CSCI 3231 Logic Circuits and Microprocessor CSCI 3232 System Software CSCI 3236 Theoretical Foundations CSCI 3432 Database Systems CSCI 4890 ...
- Bachelor of Science in Computer Science - (CEIT) | Georgia Southern ...**
ceit.georgiasouthern.edu/cs/degrees/bscs/ ▾
Aug 25, 2017 - Learn more about Georgia Southern University's bachelor's computer science degree. ... A student must have completed both MATH 2130 and CSCI 1301 (Programming Principles I) with a "C" or better and be enrolled in CSCI 1302 in order to be accepted as a major in Computer Science. Otherwise ...
- CSCI Computer Science < Georgia Southern University | A Public ...**
catalog.georgiasouthern.edu/undergraduate/course-descriptions/courses/csci/ ▾
CSCI 1302 Programming Principles II. 3 Credit Hours. 3 Lecture Hours. 0 Lab Hours. This course is a continuation of CSCI 1301. Emphasis is on advanced techniques such as recursion, regular expressions, refactoring, object oriented programming concepts and constructs, reusing components, templates/generics, ...
- CSCI 1302 : Programming Principles II - GSU - Page 1 - Course**
<https://www.coursehero.com/Schools/Georgia-Southern-University/CSCI/> ▾
Here is the best resource for homework help with CSCI 1302 : Programming Principles II at Georgia Southern University. Find CSCI1302 study guides, notes, and.
- CSCI 1302: Programming Principles II: Georgia Southern (GSU): Koofers**
www.koofers.com/.../CSCI-CSCI-Computer-Science ▾
CSCI 1302 - Programming Principles II at Georgia Southern University. CSCI 1302 · Teacher Ratings · Read student teacher reviews · CSCI 1302 · Past Tests & Quizzes · See previous tests and quizzes · CSCI 1302 · Grade Histories · See how teachers graded in the past · CSCI 1302 · Flashcard Decks · Created by students ...
- CSCI 4439 - Game Programming at Georgia Southern University ...**
<https://www.coursicle.com/georgiasouthern/courses/CSCI/4439/> ▾
CSCI 4439 at Georgia Southern University (GS). An introduction to game design and development including game physics, using game engines, using AI in games, creating multithreaded games, and creating networked games. Prerequisite(s): CSCI 1302 or permission of instructor. Cross Listing(s): CSCI 4439H.

Examples of Information Retrieval



Examples of Information Retrieval



Google search results for "machine learning" showing book results:

All News Videos Images Books More Settings Tools

About 1,200,000 results (0.55 seconds)

Python Machine Learning
<https://books.google.com/books?isbn=1783555149>
Sebastian Raschka - 2015 - Preview - More editions
Unlock deeper insights into Machine Learning with this vital guide to cutting-edge predictive analytics About This Book Leverage Python's most powerful open-source libraries for deep learning, data wrangling, and data visualization Learn ...

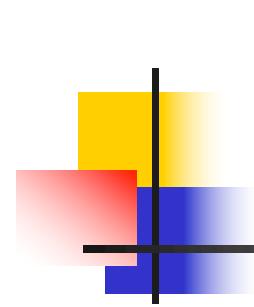
Machine Learning: An Artificial Intelligence Approach - Volume 3
<https://books.google.com/books?isbn=0080510558>
Yves Kodratoff, Ryszard S. Michalski - 2014 - Preview - More editions
Machine Learning: An Artificial Intelligence Approach, Volume III presents a sample of machine learning research representative of the period between 1986 and 1989. The book is organized into six parts.

Machine Learning: The New AI
<https://books.google.com/books?isbn=0262529513>
Ethem Alpaydin - 2016 - Preview - More editions
"A concise overview of machine learning, a type of artificial intelligence that enables computers to learn from large data sets and is the basis of numerous developments, from speech recognition to driverless cars."

Machine Learning: A Probabilistic Perspective
<https://books.google.com/books?isbn=0262018020>
Kevin P. Murphy - 2012 - Preview - More editions
The book is suitable for upper-level undergraduates with an introductory-level college math background and beginning graduate students.

Machine Learning: A Constraint-Based Approach
<https://books.google.com/books?isbn=0081006705>
Marco Gori - 2017 - Preview - More editions
This book presents a simpler unified notion of regularization, which is strictly connected with the parsimony principle, and includes many solved exercises that are classified according to the Donald Knuth ranking of difficulty, which ...

Elements of Machine Learning
<https://books.google.com/books?isbn=1558603018>
Pat Langley - 1996 - Preview
Machine learning is the computational study of algorithms that improve performance based on experience, and this book covers the basic issues of artificial intelligence.

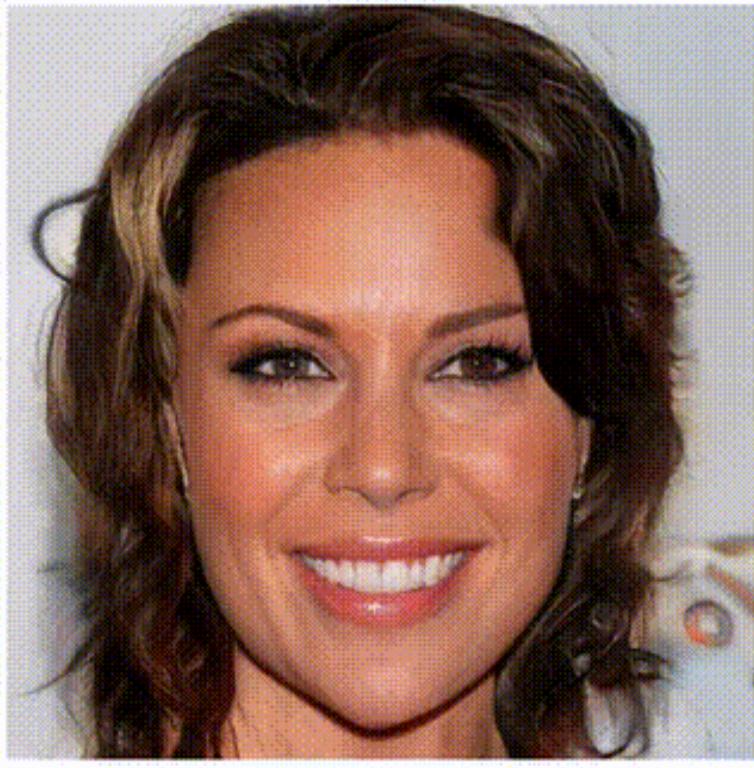


Learning algorithms are useful in many tasks

1. **Classification:** Determine which discrete category the example is
2. **Recognizing patterns:** Speech Recognition, facial identity, etc
3. **Recommender Systems:** Noisy data, commercial pay-off (e.g., Amazon, Netflix).
4. **Information retrieval:** Find documents or images with similar content
5. **Computer vision:** detection, segmentation, depth estimation, optical flow, etc

Examples of Computer Vision

INSTRUCTION: press +/- to adjust feature, toggle feature name to lock the feature



random face		
Male	Age	Skin_Tone
-	+	-
Bangs	Hairline	Bald
-	+	-
Big_Nose	Pointy_Nose	Makeup
-	+	-
Smiling	Mouth_Open	Wavy_Hair
-	+	-
Beard	Goatee	Sideburns
-	+	-
Blond_Hair	Black_Hair	Gray_Hair
-	+	-
Eyeglasses	Earrings	Necktie
-	+	-

<https://blog.insightdatascience.com/generating-custom-photo-realistic-faces-using-ai-d170b1b59255>

Examples of Computer Vision



Fig: <https://i.ytimg.com/vi/CrnLINlbFA/hqdefault.jpg>

<https://www.youtube.com/watch?v=CrnLINlbFA>

Examples of Computer Vision

A



B



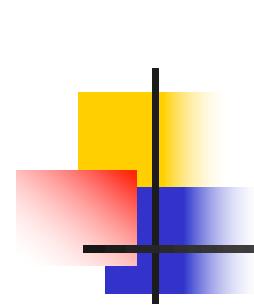
C



D



[Gatys, Ecker, Bethge. A Neural Algorithm of Artistic Style. Arxiv'15.]



Learning algorithms are useful in many tasks

1. **Classification:** Determine which discrete category the example is
2. **Recognizing patterns:** Speech Recognition, facial identity, etc
3. **Recommender Systems:** Noisy data, commercial pay-off (e.g., Amazon, Netflix).
4. **Information retrieval:** Find documents or images with similar content
5. **Computer vision:** detection, segmentation, depth estimation, optical flow, etc
6. **Robotics:** perception, planning, etc

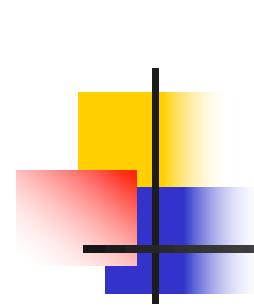
Autonomous Driving



Flying Robots

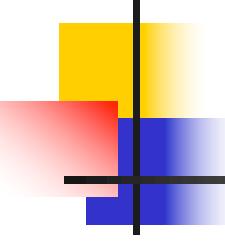


Figure: Video: <https://www.youtube.com/watch?v=YQIMGV5vtd4>



Learning algorithms are useful in many tasks

1. **Classification:** Determine which discrete category the example is
2. **Recognizing patterns:** Speech Recognition, facial identity, etc
3. **Recommender Systems:** Noisy data, commercial pay-off (e.g., Amazon, Netflix).
4. **Information retrieval:** Find documents or images with similar content
5. **Computer vision:** detection, segmentation, depth estimation, optical flow, etc
6. **Robotics:** perception, planning, etc
7. **Learning to play games**



Playing Games: Atari

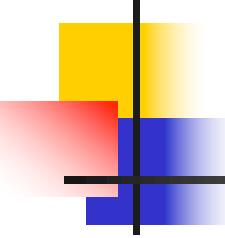


Figure: Video: <https://www.youtube.com/watch?v=V1eYniJ0Rnk>



Learning algorithms are useful in many tasks

1. **Classification:** Determine which discrete category the example is
2. **Recognizing patterns:** Speech Recognition, facial identity, etc
3. **Recommender Systems:** Noisy data, commercial pay-off (e.g., Amazon, Netflix).
4. **Information retrieval:** Find documents or images with similar content
5. **Computer vision:** detection, segmentation, depth estimation, optical flow, etc
6. **Robotics:** perception, planning, etc
7. **Learning to play games**
8. **Recognizing anomalies:** Unusual sequences of credit card transactions, panic situation at an airport
9. **Spam filtering, fraud detection:** The enemy adapts so we must adapt too
10. **Many more!**



Types of Learning

- **Supervised (inductive) learning**

- Training data includes desired outputs

- **Unsupervised learning**

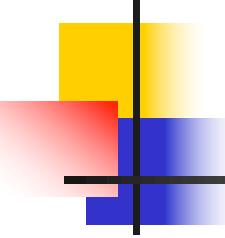
- Training data does not include desired outputs

- **Semi-supervised learning**

- Training data includes a few desired outputs

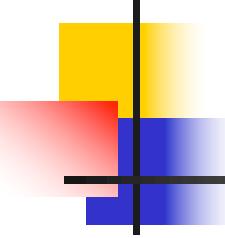
- **Reinforcement learning**

- Rewards from sequence of actions



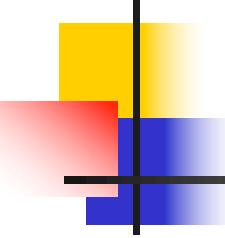
Supervised Learning

- **Given** examples of a function ($X, F(X)$)
- **Predict** function $F(X)$ for new examples X
 - **Discrete $F(X)$** : Classification (1-of-N output (speech recognition, object recognition, medical diagnosis))
 - **Continuous $F(X)$** : Regression (real-valued output (predicting market prices, customer rating))
 - $F(X) = \text{Probability}(X)$: Probability estimation



Unsupervised Learning

- Create an internal representation of the input, capturing regularities/structure in data
- Example:
 - Clustering: Discover groups of similar inputs (documents, images, etc)



What We'll Cover

- **Supervised learning**

- Decision tree induction
- Instance-based learning
- Linear Regression
- Bayesian learning
- Logistic Regression
- Neural networks

- **Unsupervised learning**

- Clustering
- Dimensionality reduction